

DISULFIND: a disulfide bonding state and cysteine connectivity prediction server

Alessio Ceroni¹, Andrea Passerini¹, Alessandro Vullo² and Paolo Frasconi^{1,*}

¹Machine Learning and Neural Networks Group, Università degli Studi di Firenze, Dipartimento di Sistemi e Informatica, Via di Santa Marta 3, 50139 Firenze, Italy and ²School of Computer Science and Informatics, University College Dublin, Belfield, Dublin 4, Ireland

Received March 14, 2006; Revised March 29, 2006; Accepted March 31, 2006

ABSTRACT

DISULFIND is a server for predicting the disulfide bonding state of cysteines and their disulfide connectivity starting from sequence alone. Optionally, disulfide connectivity can be predicted from sequence and a bonding state assignment given as input. The output is a simple visualization of the assigned bonding state (with confidence degrees) and the most likely connectivity patterns. The server is available at <http://disulfind.dsi.unifi.it/>.

INTRODUCTION

Disulfide bridges play a major role in the stabilization of the folding process and, consequently, in studies related to structural and functional properties of specific proteins. In addition, knowledge about the disulfide bonding state of cysteines may help the experimental structure determination process and may be useful in other genomic annotation tasks.

DISULFIND uses a combination of machine learning algorithms to predict intrachain bridges from sequence alone. Similar to many other tools of this kind, it solves the prediction problem in two steps. First, the disulfide bonding state of each cysteine is predicted by a binary classifier; second, cysteines that are known to participate in the formation of bridges are paired to obtain a connectivity pattern.

RELATED WORKS

Early work on bonding state employed representations based on local-window multiple alignment profiles and neural networks for discrimination (1,2). Mucchielli-Giorgi *et al.* (3) introduced the idea of adding a global descriptor to improve prediction accuracy. Ceroni *et al.* (4) proposed a method based on a combination of string and vector kernels in conjunction with support vector machines (SVMs). Song *et al.* (5)

applied a linear discriminant using dipeptides as features. Martelli *et al.* (6) suggested the use of hidden Markov models to refine local predictions obtained via neural networks. SVMs are also used in the method presented in (7).

Prediction of connectivity patterns was pioneered in (8) with a method based on weighted graph matching, implemented in the prediction server DCON. Vullo and Frasconi (9) introduced the use of multiple alignment profiles by means of recursive neural networks (RNNs). In this approach, (that still underpins DISULFIND) a global score is assigned to an entire connectivity pattern. In the DAG RNN approach described in (7,10), the probability for a disulfide bond is computed for each pair of cysteines. The associated Dipro server (which also predicts bonding state) is described in (11). Taskar *et al.* (12) formulated disulfide connectivity as a structured-output prediction problem and solved it using a generalized large-margin machine. Ferrè and Clote (13) proposed a feedforward neural-network architecture with hidden units associated with cysteine pairs and inputs encoding secondary structure; the method is behind the prediction server DiANNA (14). Zhao *et al.* (15) confirmed that the profile of distances between bonded cysteines is an important feature for prediction of connectivity patterns. This idea has been further exploited in conjunction with SVMs to develop the method behind the prediction server PreCys (16). Finally, CysView (17) is a server that predict patterns by comparison of a query sequence to annotated data bases.

MATERIALS AND METHODS

Multiple alignment profiles

Prediction of protein structural properties is typically more accurate when incorporating evolutionary information encoded in multiple alignment profiles. Profiles are used in DISULFIND both in bonding state and connectivity prediction. They are calculated by using one iteration of the

*To whom correspondence should be addressed. Tel: +39 0554796362; Fax: +39 0554796363; Email: p-f@dsi.unifi.it

The authors wish it to be known that, in their opinion, the first three authors should be regarded as joint first authors.

© The Author 2006. Published by Oxford University Press. All rights reserved.

The online version of this article has been published under an open access model. Users are entitled to use, reproduce, disseminate, or display the open access version of this article for non-commercial purposes provided that: the original authorship is properly and fully attributed; the Journal and Oxford University Press are attributed as the original place of publication with the correct citation details given; if an article is subsequently reproduced or disseminated not in its entirety but only in part or as a derivative work this must be clearly indicated. For commercial re-use, please contact journals.permissions@oxfordjournals.org

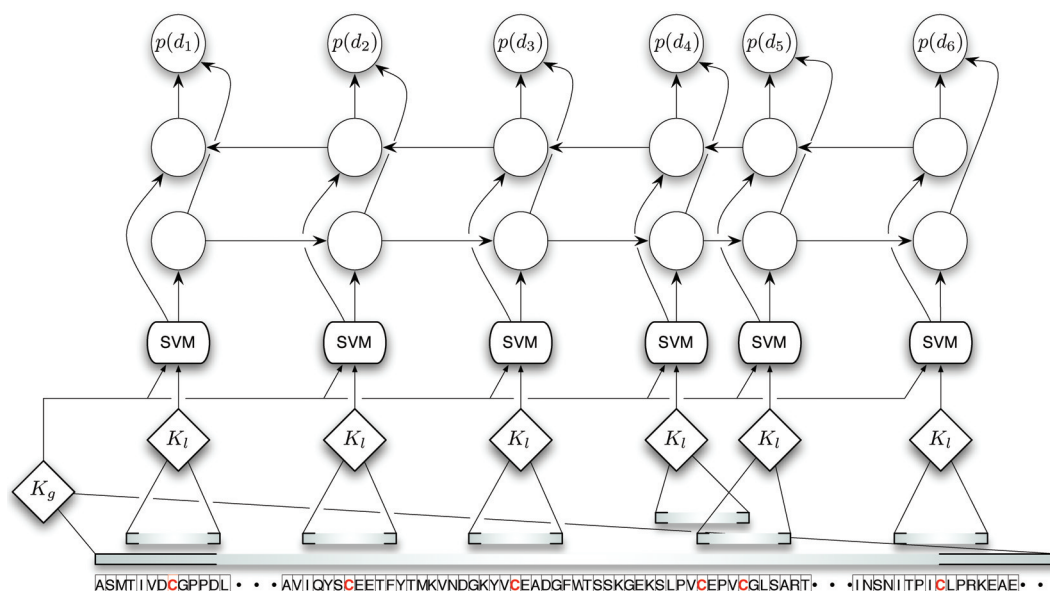


Figure 1. Architecture of the bonding state predictor. The lower level provides independent cysteine predictions based on a local kernel k_l on local attributes, and a global kernel k_g on the entire sequence. The upper level is a BRNN (represented here schematically by its graphical model) that outputs a disulfide-bonding probability $p(d_i)$ for each cysteine, based on all SVM predictions.

PSI-BLAST program run on Swiss-Prot and TrEMBL using the BLOSUM62 matrix and an E -value cutoff of 0.005.

Prediction of disulfide bonding state

DISULFIND employs an SVM binary classifier to predict the bonding state of each cysteine, followed by a refinement stage that classifies all the cysteines in a chain in a collective fashion (18), that is, by deciding the overall bonding state assignment of an entire chain rather than making several independent predictions (one for each cysteine). The overall architecture is shown in Figure 1. The SVM receives as input both local and global features [see also (3)]. Local features consist of a window of position specific conservations derived from multiple alignment, centered around the target residue. Global features (amino acid composition, chain length, number of cysteines and average cysteine conservation) provide information about the bonding class of the entire chain (all cysteines bonded, none or mix), which is strongly correlated with the subcellular compartment where the protein resides (reducing versus oxidizing environments).

The refinement stage is motivated by the observation that single cysteines are not independently sampled. Linkage occurs between pairs forming a disulfide bridge but also among sets of cysteines that coordinate a metal ion. A second source of linkage is due to the fact that bonding state is very often a global property of the protein chain and not a local property of individual cysteines (2,3). The effects of correlation are mitigated in two ways. First, we trained a bidirectional recurrent neural network (BRNN) (19) to predict a globally correct sequence of bonding state assignments, given a (possibly incorrect) sequence of locally calculated predictions. At each cysteine position i , the BRNN output is computed using the logistic function and can be therefore interpreted as the conditional probability $p(d_i)$ that the




Figure 2. Finite state automaton used in the final stage of bonding state prediction.

cysteine is disulfide-bonded given the input sequence. A position-specific prediction confidence is then defined as

$$c_i = 2(\max\{p(d_i), 1 - p(d_i)\} - 0.5). \quad 1$$

Second, we enforce the number of bonded cysteines to be even (interchain bridges are ignored) using a finite state automaton (shown in Figure 2). Given the sequence of bonding state probabilities (computed by the BRNN), the most likely sequence of bonding states is obtained by running a Viterbi algorithm. Similar ideas (but using a hidden Markov models rather than an automaton) were presented in (6).

Prediction of disulfide connectivity

We assume in this subsection that disulfide-bonding state of cysteines is given (either entered manually by the user or predicted using the method described above). The method used in DISULFIND is fully detailed in (9) and briefly summarized here.

A connectivity pattern can be conveniently represented as an undirected graph whose vertices are cysteines and edges

Email Address

Query Name

Amino Acid Sequence, single letter code (*)

Predict Options:

- ☐ Predict connectivity from user specified bonding state
- ☒ Predict bonding state + connectivity pattern

Output Options:

- ☐ Batch: send output to email address
- ☒ Interactive: receive output on browser

Number of alternative patterns to be returned

Before submitting your first query please read the following note

Figure 3. Screenshot of the DISULFIND input form.

are disulfide bridges. The problem thus consists of mapping an input sequence with annotated cysteines into an output graphs representing disulfide connectivity. This structured output prediction problem can be cast in the traditional supervised learning setting by introducing a regression problem defined as follows. The input is formed by the annotated sequence and a candidate connectivity pattern. The target is a real valued score, defined as the fraction of correctly assigned bridges. During training the target score is known and we use it to train a recursive neural network in regression mode. Prediction is carried out by running the trained network on all possible connectivity patterns and choosing the one yielding maximum score. The number of possible disulfide patterns connecting $2B$ cysteines is $(2B-1)!!$ where the double factorial $n!!$ is defined as the product of all odd integers that are less or equal to n .

In order to limit computational efforts, DISULFIND can assign at most five disulfide bridges (in this case the number of candidates to be evaluated is 945). Two remarks are relevant to this limitation. First, chains with more than five bridges are rare (no more than 10% of the Swiss-Prot chains annotated with disulfide bridges). Second, the prediction accuracy is already low for chains having five bridges because of a limited number of available training examples; hence prediction of patterns with six or more bridges would be very inaccurate.

IMPLEMENTATION

DISULFIND is available both as a standalone service at <http://disulfind.dsi.unifi.it/> and as part of PredictProtein (20).

The current version (DISULFIND 1.1, released in February 2006), incorporates some improvements in the presentation interface.

Interface

The input to the predictor is entered via an HTTP form using the SEND method. The main fields (see Figure 3) are the following.

Email address The address where results will be sent if the email output option is selected.

Query name An optional field that allows to label the sequence with an user provided identifier.

Amino acid sequence The protein sequence using standard amino acid one-letter codes. Spaces and newlines are automatically stripped.

Predict options In its normal behavior, DISULFIND predicts both bonding state and connectivity. If the bonding state is known in advance, users may check the corresponding option in the user interface and after the form is submitted they will be presented a screen where the bonding state of each cysteine can be manually assigned. In this case only predicted connectivity will be returned.

Output options There are two possible output operation modes. In email mode, after the form is submitted, a job is scheduled in the server and results are returned in ASCII format to the indicated email address. In browser mode, results are returned to the HTTP client (see Figure 4).

Alternatives By default DISULFIND only returns the most likely connectivity pattern. By setting the number of alternatives to an integer k in the range (1,3), the k best ranking patterns will be returned.

Table 2. Leave-one-out validation results of disulphide connectivity prediction

Number of bridges	Number of chains	Q_p	Q_c
2	156	75.0	75.0
3	146	46.6	55.7
4	99	50.5	63.4
5	45	17.8	42.7
All	446	54.5	60.2

Concerning disulfide connectivity, leave-one-out estimates of prediction accuracy on a set of 446 Swiss-Prot Sequences (9) are reported in Table 2 [note that results reported in (9) were based on a 4-fold cross validation]. Q_p is the fraction of correctly assigned patterns, while Q_c is the fraction of correctly predicted bridges. If multiple alternative are selected, the probability that a correct pattern is included increases. Results obtained considering the top $k = 3$ configurations are $Q_p = 66.3$, $Q_c = 69.5$.

Statistics

DISULFIND has served a total of over 7000 tasks from almost 50 national domains since April 2003 and is currently serving an average of 60 queries per week. Hundreds of queries per month have been served via PredictProtein since July 2004.

ACKNOWLEDGEMENTS

The authors would like to thank Rita Casadio and Piero Fariselli for useful discussions and the anonymous reviewers for their suggestions. This research is supported by EU STREP APriL II (contract no. FP6-508861) and EU NoE BIOPATTERN (contract no. FP6-508803). The work of A.V. is supported by an Embark Fellowship from the Irish Research Council for Science, Engineering and Technology. The Open Access publication charges for this article were waived by Oxford University Press.

Conflict of interest statement. None declared.

REFERENCES

- Fariselli, P., Riccobelli, P. and Casadio, R. (1999) Role of evolutionary information in predicting the disulfide-bonding state of cysteine in proteins. *Proteins*, **36**, 340–346.
- Fiser, A. and Simon, I. (2000) Predicting the oxidation state of cysteines by multiple sequence alignment. *Bioinformatics*, **16**, 251–256.
- Mucchielli-Giorgi, M.H., Hazout, S. and Tuffery, P. (2000) Predicting the disulfide bonding state of cysteines using protein descriptors. *Proteins*, **46**, 243–249.
- Ceroni, A., Frasconi, P., Passerini, A. and Vullo, A. (2003) Predicting the disulfide bonding state of cysteines with combinations of kernel machines. *J. VLSI Signal Processing*, **35**, 287–295.
- Song, J.-N., Wang, M.-L., Li, W.-J. and Xu, W.-B. (2004) Prediction of the disulfide-bonding state of cysteines in proteins based on dipeptide composition. *Biochem. Biophys. Res. Commun.*, **318**, 142–147.
- Martelli, P.L., Fariselli, P. and Casadio, R. (2004) Prediction of disulfidebonded cysteines in proteomes with a hidden neural network. *Proteomics*, **4**, 1665–1671.
- Cheng, J., Saigo, H. and Baldi, P. (2006) Large-scale prediction of disulphide bridges using kernel methods, two-dimensional recursive neural networks, and weighted graph matching. *Proteins*, **62**, 617–629.
- Fariselli, P. and Casadio, R. (2001) Prediction of disulfide connectivity in proteins. *Bioinformatics*, **17**, 957–964.
- Vullo, A. and Frasconi, P. (2004) Disulfide connectivity prediction using recursive neural networks and evolutionary information. *Bioinformatics*, **20**, 653–659.
- Baldi, P., Cheng, J. and Vullo, A. (2005) Large-scale prediction of disulphide bond connectivity. In Saul, L.K., Weiss, Y. and Bottou, L. (eds), *Advances in Neural Information Processing Systems 17*. MIT Press, Cambridge, MA, pp. 97–104.
- Cheng, J., Randall, A.Z., Sweredoski, M.J. and Baldi, P. (2005) SCRATCH: a protein structure and structural feature prediction server. *Nucleic Acids Res.*, **33**, W72–W76.
- Taskar, B., Chatalbashev, V., Koller, D. and Guestrin, C. (2005) Learning structured prediction models: a large margin approach. In *Proceedings of the Twenty Second International Conference on Machine Learning (ICML05)*.
- Ferrè, F. and Clote, P. (2005) Disulfide connectivity prediction using secondary structure information and diresidue frequencies. *Bioinformatics*, **21**, 2336–2346.
- Ferrè, F. and Clote, P. (2005) DiANNA: a web server for disulfide connectivity prediction. *Nucleic Acids Res.*, **33**, W230–W232.
- Zhao, E., Liu, H.-L., Tsai, C.-H., Tsai, H.-K., Chan, C.-L. and Kao, C.-Y. (2005) Cysteine separations profiles on protein sequences infer disulfide connectivity. *Bioinformatics*, **21**, 1415–1420.
- Tsai, C.-H., Chen, B.-J., Chan, C.-H., Liu, H.-L. and Kao, C.-Y. (2005) Improving disulfide connectivity prediction with sequential distance between oxidized cysteines. *Bioinformatics*, **21**, 4416–4419.
- Lenffer, J., Lai, P., El Mejaber, W., Khan, A.M., Koh, J.L.Y., Tan, P.T.J., Seah, S.H. and Brusic, V. (2004) CysView: protein classification based on cysteine pairing patterns. *Nucleic Acids Res.*, **32**, W350–W355.
- Getoor, L., Friedman, N., Koller, D. and Taskar, B. (2001) Learning probabilistic models of relational structure. In *Proceedings 18th International Conf. on Machine Learning*. Morgan Kaufmann, San Francisco, CA, pp. 170–177.
- Baldi, P., Brunak, S., Frasconi, P., Soda, G. and Pollastri, G. (1999) Exploiting the past and the future in protein secondary structure prediction. *Bioinformatics*, **15**, 937–946.
- Rost, B., Yachdav, G. and Liu, J. (2004) The PredictProtein server. *Nucleic Acids Res.*, **32**, W321–W326.